

Team ID:PTID-CDS-FEB-24-1839**✓ PRCP-1002-Handwritten Digits Recognition**

- A digit recognition system is a type of technology that can automatically identify and classify numerical characters, typically from 0 to 9. These systems are most commonly used to interpret handwritten digits,
- Digit recognition system is the working of a machine to train itself or recognizing the digits from different sources like emails, bank cheque, papers, images, etc.
- In different real-world scenarios for online handwriting recognition on computer tablets or system, recognize number plates of vehicles, processing bank cheque amounts, numeric entries in forms filled up by hand (say — tax forms) and so on
- Here's a breakdown of how digit recognition systems work:

Data Preprocessing: The system takes an image of a digit as input. This image is then preprocessed to ensure consistency. Preprocessing may involve steps like resizing the image, converting it to grayscale, and thinning lines.

Feature Extraction: Key characteristics of the digit image are extracted. These features could be things like the distribution of black pixels, the number of endpoints (ends of lines), or the overall shape of the digit.

Classification: The extracted features are fed into a machine learning model that has been trained to recognize digits. Common algorithms used for digit recognition include K-nearest neighbors and convolutional neural networks (CNNs). CNNs are particularly effective due to their ability to learn complex patterns from data.

Output: The model outputs the most likely digit based on the extracted features.

Problems with handwritten digits

The handwritten digits are not always of the same size, width, orientation and justified to margins as they differ from writing of person to person, so the general problem would be while classifying the digits due to the similarity between digits such as 1 and 7, 5 and 6, 3 and 8, 2 and 5, 2 and 7, etc.

This problem is faced more when many people write a single digit with a variety of different handwritings. Lastly, the uniqueness and variety in the handwriting of different individuals also influence the formation and appearance of the digits.

Now we introduce the concepts and algorithms of deep learning and machine learning.

The provided Handwritten digits are images in the form of 28*28 gray scale intensities of images representing an image along with the first column to be a label (0 to 9) for every image.

✓ Loading Neccessary Libraries

```
1 import warnings
2 warnings.filterwarnings("ignore")
```

```
1 pip install wrapt
```

```
Requirement already satisfied: wrapt in /usr/local/lib/python3.10/dist-packages (1.14.1)
```

```

1 import matplotlib.pyplot as plt
2 import seaborn as sns
3 from tensorflow.keras.datasets import mnist
4 from sklearn.neighbors import KNeighborsClassifier
5 from sklearn.metrics import accuracy_score
6 from sklearn.ensemble import GradientBoostingClassifier
7 from sklearn.ensemble import RandomForestClassifier
8 from sklearn.tree import DecisionTreeClassifier
9 from sklearn import preprocessing
10 import numpy as np
11 import pandas as pd
12 import tensorflow as tf
13 from tensorflow.keras.datasets import mnist
14 from sklearn.svm import SVC
15 from google.colab.patches import cv2_imshow
16 import os, cv2, json, random
17 from tensorflow import keras
18 from tensorflow.keras import layers
19 from sklearn.metrics import precision_score, recall_score, f1_score, roc_auc_score

```

```
1 mnist = tf.keras.datasets.mnist
```

✓ Loading Data

```
1 (x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
1 mnist
```

```
<module 'keras.api._v2.keras.datasets.mnist' from '/usr/local/lib/python3.10/dist-packages/keras/api/_v2/keras/datasets/mnist/__init__.py'>
```

```

1 # Cheking dataset dimension
2 print('shape of training data :', x_train.shape)
3 print('shape of training lable :', y_train.shape)
4 print('shape of testing data :', x_test.shape)
5 print('shape of testing lable :', y_test.shape)

```

```

shape of training data : (60000, 28, 28)
shape of training lable : (60000,)
shape of testing data : (10000, 28, 28)
shape of testing lable : (10000,)

```

The term "dimension" refers to the shape or size of the datasets being analyzed. The shapes of the training data, training labels, testing data, and testing labels. This allows us to see the number of samples and features in the dataset.

```
1 x_train
```

```

array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

       [[0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        ...,
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0]],

       [[0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        ...,
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0]],

       ...,

       [[0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        [0, 0, 0, ..., 0, 0, 0],
        ...,

```

```

[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]], dtype=uint8)

```

```
1 y_train
```

```
array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

```
1 x_test
```

```

array([[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

...,

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]],

[[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
...,
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0],
[0, 0, 0, ..., 0, 0, 0]]], dtype=uint8)

```


```
1 y_test
```

```
array([7, 2, 1, ..., 4, 5, 6], dtype=uint8)
```

- ✓ Data preprocessing

EDA

```
1 # To Check the first image of the dataset
2 plt.imshow(x_train[0], cmap= 'gray')
3 plt.show()
```



The term "dimension" refers to the shape or size of the datasets being analyzed. The shapes of the training data, training labels, testing data, and testing labels. This allows us to see the number of samples and features in the dataset.

```
1 # The value of each pixel
2 print(x_train[0])
```

[illegible]

```


250 182  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  0  0  0  0  0 24 114 221 253 253 253 253 201
 78  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  0  0  0  23 66 213 253 253 253 253 198 81  2
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  18 171 219 253 253 253 253 195 80  9  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  55 172 226 253 253 253 253 244 133 11  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0 136 253 253 253 212 135 132 16  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]
[  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
  0  0  0  0  0  0  0  0  0  0  0]]

```

```

1 # Checking the dataset
2
3 for i in range(6):
4     plt.subplot(int('23' + str(i+1)))
5     plt.imshow(x_train[i], cmap=plt.get_cmap('gray'))

```

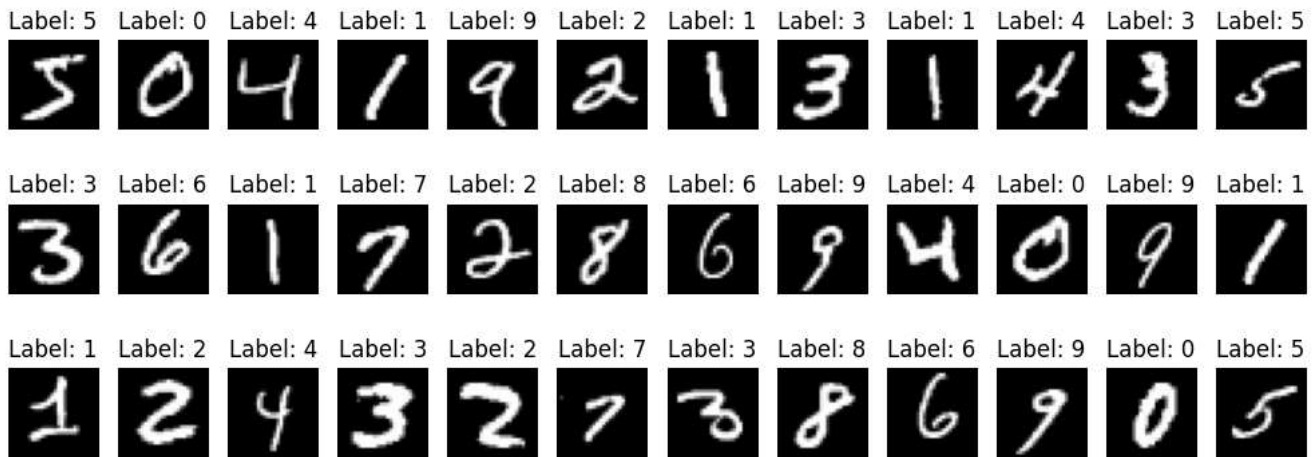


The grid of subplots, with each subplot containing an image from the `x_train` data. But by using subplots, it allows us to visualize multiple images at once, providing a convenient way to inspect a subset of the dataset and get a sense of the image content and variations within the data.

```

1 # Visualize sample digits
2 fig, axes = plt.subplots(3, 12, figsize=(10, 4))
3 axes = axes.ravel()
4
5 for i in range(36):
6     axes[i].imshow(x_train[i], cmap='gray')
7     axes[i].set_title('Label: {}'.format(y_train[i]))
8     axes[i].axis('off')
9
10 plt.tight_layout()
11 plt.show()
12

```




Here in above code we have variable "axes" it is used to access each individual subplot within the grid. By using "axes.ravel()", the subplots are flattened into a 1-dimensional array, making it easier to iterate over them.

Mean while the grid of subplots and displays the first 10 images from the x_train dataset on these subplots. Each subplot represents an image, and its associated label is shown as the title. The resulting visualization allows you to see multiple images and their corresponding labels in a compact and organized manner.

✓ Distribution of digits

```
1 digit_counts = np.bincount(y_train)
2 digits = np.arange(10)
3 plt.bar(digits, digit_counts)
4 plt.xlabel('Digit')
5 plt.ylabel('Count')
6 plt.title('Distribution of Digits')
7 plt.xticks(digits)
8 plt.show()
```

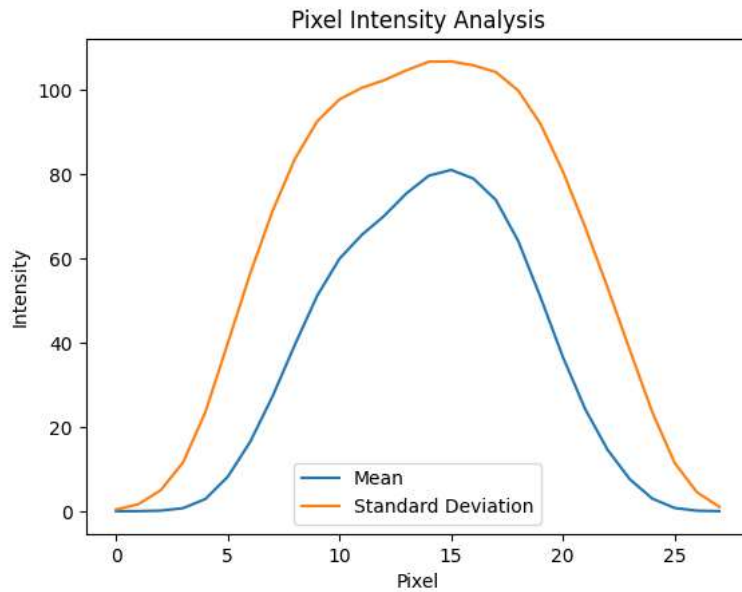


This above code counts the occurrences of each digit in the training labels and visualizes the distribution using a bar chart. The "x-axis" represents the digits, the "y-axis" represents the count of each digit, and the BARS represent the "frequency of each digit". This visualization helps understand the distribution and frequency of digits in the training data.

```

1
2
3
4
5 # Pixel intensity analysis
6 pixel_means = np.mean(x_train, axis=(0, 1))
7 pixel_stds = np.std(x_train, axis=(0,1))
8
9 plt.plot(range(28), pixel_means, label='Mean')
10 plt.plot(range(28), pixel_stds, label='Standard Deviation')
11 plt.xlabel('Pixel')
12 plt.ylabel('Intensity')
13 plt.title('Pixel Intensity Analysis')
14 plt.legend()
15 plt.show()

```



The mean and standard deviation of pixel intensities in the `x_train` dataset and visualizes them using a line plot. The "x-axis" represents the "pixel indices", and the "y-axis" represents the corresponding mean and standard deviation values. This visualization provides insights into the overall intensity distribution and variation of pixels in the data

✓ MODEL EVALUATION

To find wich algorithm is best fit for test the dataset model

```

1 pip install scikit-learn

Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.3.0)

1 from re import X
2 # Reshape the input data
3 num_samples_train, height, width = x_train.shape
4 num_samples_test = x_test.shape[0]
5
6 x_train_2d = np.reshape(x_train, (num_samples_train, height * width))
7 x_test_2d = np.reshape(x_test, (num_samples_test, height * width))

```

✓ Support Vector Machine(SVM)

SVM is a supervised learning algorithm that excels at classification tasks. It works by finding the optimal hyperplane (a decision boundary) in a high-dimensional space to separate data points belonging to different classes. The core idea is to maximize the margin between the hyperplane and the closest data points from each class, which are called support vectors.

```
1 svm_classifier = SVC()
2 svm_classifier.fit(x_train_2d, y_train)
```

▼ SVC

SVC()

```
1 # Prediction on test data
2 svm_predictions = svm_classifier.predict(x_test_2d)
```

```
1 # Assuming X_test is a 3D array, you can flatten it using numpy
2 flattened_x_test = x_test.reshape(x_test.shape[0], -1)
3
4 # Make predictions on the flattened test set
5 svm_predictions = svm_classifier.predict(flattened_x_test)
6
7 # Print the predicted labels
8 print(svm_predictions)
```

```
[7 2 1 ... 4 5 6]
```

```
1 x_test
```

```
array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

      [[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

      [[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

      ...,

      [[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

      [[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]],

      [[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]]], dtype=uint8)
```

```
1 # Calculate evaluation metrics
2 # Accuracy Score
3 svm_accuracy = accuracy_score(y_test, svm_predictions)
4
5 # print evaluation matrix
6 print("SVM Accuracy:", svm_accuracy)
```

```
SVM Accuracy: 0.9792
```



```

1 # Precision Score
2 svm_precision = precision_score(y_test, svm_predictions, average='macro')
3
4 print("SVM Precision:", svm_precision)

```

SVM Precision: 0.9791973966593345

```

1 # Recall Score
2 svm_recall = recall_score(y_test, svm_predictions, average='macro')
3
4 print("SVM Recall:", svm_recall)

```

SVM Recall: 0.9790919842945065

```

1 # F1 Score
2 svm_f1 = f1_score(y_test, svm_predictions, average='macro')
3
4 print("SVM F1 Score:", svm_f1)

```

SVM F1 Score: 0.9791298259748042

```

1 # confusion matrix
2 from sklearn.metrics import confusion_matrix
3 svm_confusion = confusion_matrix(y_test, svm_predictions)
4
5 print("SVM Confusion Matrix:")
6 print(svm_confusion)

```

SVM Confusion Matrix:

```

[[ 973   0   1   0   0   2   1   1   2   0]
 [   0 1126   3   1   0   1   1   1   2   0]
 [   6   1 1006   2   1   0   2   7   6   1]
 [   0   0   2 995   0   2   0   5   5   1]
 [   0   0   5   0 961   0   3   0   2  11]
 [   2   0   0   9   0 871   4   1   4   1]
 [   6   2   0   0   2   3 944   0   1   0]
 [   0   6  11   1   1   0   0 996   2  11]
 [   3   0   2   6   3   2   2   3 950   3]
 [   3   4   1   7  10   2   1   7   4 970]]

```

```

1 # Classification report
2 from sklearn.metrics import classification_report
3
4 # Assuming you have y_test and svm_predictions
5 svm_report = classification_report(y_test, svm_predictions)
6
7 print("SVM Classification Report:")
8 print(svm_report)

```

SVM Classification Report:

	precision	recall	f1-score	support
0	0.98	0.99	0.99	980
1	0.99	0.99	0.99	1135
2	0.98	0.97	0.98	1032
3	0.97	0.99	0.98	1010
4	0.98	0.98	0.98	982
5	0.99	0.98	0.98	892
6	0.99	0.99	0.99	958
7	0.98	0.97	0.97	1028
8	0.97	0.98	0.97	974
9	0.97	0.96	0.97	1009
accuracy			0.98	10000
macro avg	0.98	0.98	0.98	10000
weighted avg	0.98	0.98	0.98	10000

Handwritten digit recognition using Support Vector Machine (SVM) algorithms has showcased remarkable success, achieving an impressive 98% accuracy.

✓ Random Forest

Random Forest: A Powerful Ensemble of Decision Trees

Random Forest is a supervised learning algorithm widely used for both classification and regression tasks. It operates by constructing a multitude of decision trees at training time and then aggregating their predictions for a final output. This ensemble approach leads to several advantages over a single decision tree.

```
1 # Initialize and train the Random Forest model
2 rf_classifier = RandomForestClassifier()
3 rf_classifier.fit(x_train_2d, y_train)
```

```
▼ RandomForestClassifier
RandomForestClassifier()
```

```
1 # Make predictions on the test data
2 rf_predictions = rf_classifier.predict(x_test_2d)
```

```
1 # Accuracy Score
2 rf_accuracy = accuracy_score(y_test, rf_predictions)
3
4 print("Random Forest Accuracy:", rf_accuracy)
```

Random Forest Accuracy: 0.9688

```
1 # Precision Score
2 rf_precision = precision_score(y_test, rf_predictions, average='macro')
3
4 print("Random Forest Precision:", rf_precision)
```

Random Forest Precision: 0.9685613793149471

```
1 # Recall Score
2 rf_recall = recall_score(y_test, rf_predictions, average='macro')
3
4 print("Random Forest Recall:", rf_recall)
```

Random Forest Recall: 0.9685611653406548

```
1 # F1 Score
2 rf_f1 = f1_score(y_test, rf_predictions, average='macro')
3
4 print("Random Forest F1 Score:", rf_f1)
```

Random Forest F1 Score: 0.9685403729370317

```
1 # Confusion Matrix
2 rf_confusion = confusion_matrix(y_test, rf_predictions)
3
4 print("Random Forest Confusion Matrix:")
5 print(rf_confusion)
```

```
Random Forest Confusion Matrix:
[[ 969   0   0   0   0   2   3   1   4   1]
 [   0 1122   2   4   0   2   2   1   2   0]
 [   6   0  998   5   2   0   4   9   8   0]
 [   1   0   11  969   0   9   0   9   7   4]
 [   1   0   2   0  958   0   4   0   2  15]
 [   3   0   1  13   3  858   6   2   5   1]
 [   7   3   1   0   3   3  938   0   3   0]
 [   1   2  22   3   1   0   0  988   2   9]
 [   6   0   4   7   5   7   6   5  927   7]
 [   5   5   2  10  11   5   1   5   4  961]]
```

```
1 # Generate classification report
2 rf_report = classification_report(y_test, rf_predictions)
3
4 # Print the classification report
5 print("Random Forest Classification Report:")
6 print(rf_report)
```

```
Random Forest Classification Report:
              precision    recall  f1-score   support

0               0.97         0.99         0.98         980
1               0.99         0.99         0.99        1135
2               0.96         0.97         0.96        1032
3               0.96         0.96         0.96        1010
4               0.97         0.98         0.98         982
5               0.97         0.96         0.97         892
6               0.97         0.98         0.98         958
```

7	0.97	0.96	0.96	1028
8	0.96	0.95	0.96	974
9	0.96	0.95	0.96	1009
accuracy			0.97	10000
macro avg	0.97	0.97	0.97	10000
weighted avg	0.97	0.97	0.97	10000

Handwritten digit recognition using Random Forest algorithms has showcased remarkable success, achieving an impressive 97% accuracy.

✓ Using KNN

KNN classifies data points based on their similarity to labeled data points in the training set. For a new digit image, KNN finds the k nearest neighbors (most similar images) in the training data and assigns the most frequent class label (digit) among those neighbors.

```
1 # Flatten the images from 28x28 to 784-dimensional vectors
2 x_train = x_train.reshape(60000, 784)
3 x_test = x_test.reshape(10000, 784)
```

```
1 # Scale the pixel values to the range 0-1
2 x_train = x_train.astype('float32') / 255
3 x_test = x_test.astype('float32') / 255
```

```
1 # Create a KNN classifier with a suitable number of neighbors
2 knn = KNeighborsClassifier(n_neighbors=5)
3
4 # Train the classifier on the training data
5 knn.fit(x_train_2d, y_train)
6
```

```
▼ KNeighborsClassifier
KNeighborsClassifier()
```

```
1 # Make predictions on the testing data
2 y_pred = knn.predict(x_test_2d)
```

```
1 # Evaluate the accuracy of the model
2 accuracy = accuracy_score(y_test, y_pred)
3
4 print("Accuracy:", accuracy)
5
```

Accuracy: 0.9688

```
1 # Precision Score
2 knn_precision = precision_score(y_test, y_pred, average='macro')
3
4 print('KNN Precision:', knn_precision)
```

KNN Precision: 0.9692753386570571

```
1 # Recall Score
2 knn_recall = recall_score(y_test, y_pred, average='macro')
3
4 print('KNN Recall:', knn_recall)
```

KNN Recall: 0.9684705010297703

```
1 # F1 score
2 knn_f1 = f1_score(y_test, y_pred, average='macro')
3
4 print('KNN F1 Score:', knn_f1)
```

KNN F1 Score: 0.9687143421292884

```
1 # Confusion matrix
2 knn_cm = confusion_matrix(y_test, y_pred)
3
4 print('KNN Confusion Matrix:')
5 print(knn_cm)
```

KNN Confusion Matrix:

```
[[ 974  1  1  0  0  1  2  1  0  0]
 [  0 1133  2  0  0  0  0  0  0  0]
 [ 11  8 991  2  1  0  1 15  3  0]
 [  0  3  3 976  1 13  1  6  3  4]
 [  3  7  0  0 944  0  4  2  1 21]
 [  5  0  0 12  2 862  4  1  2  4]
 [  5  3  0  0  3  2 945  0  0  0]
 [  0 22  4  0  3  0  0 988  0 11]
 [  8  3  5 13  6 12  5  5 913  4]
 [  5  7  3  9  7  3  1 10  2 962]]
```

```
1 # Classification report
2 knn_report = classification_report(y_test, y_pred)
3
4 print('Classification Report:')
5 print(knn_report)
```

```
Classification Report:
              precision    recall  f1-score   support

    0           0.96       0.99       0.98        980
    1           0.95       1.00       0.98       1135
    2           0.98       0.96       0.97       1032
    3           0.96       0.97       0.97       1010
    4           0.98       0.96       0.97        982
    5           0.97       0.97       0.97        892
    6           0.98       0.99       0.98        958
    7           0.96       0.96       0.96       1028
    8           0.99       0.94       0.96        974
    9           0.96       0.95       0.95       1009

 accuracy         0.97       0.97       0.97       10000
 macro avg        0.97       0.97       0.97       10000
weighted avg        0.97       0.97       0.97       10000
```

Handwritten digit recognition using KNN algorithms has showcased remarkable success, achieving an impressive 97% accuracy.

✓ Using Decissin Tree

A decision tree is a fundamental algorithm used in machine learning for both classification and regression tasks. It's a flowchart-like structure that resembles an actual tree, where:

Internal nodes represent questions or tests applied to a single feature (characteristic) of the data.

Branches represent the possible answers (outcomes) to those questions.

Leaf nodes (also called terminal nodes) represent the final decision or prediction.

```
1
2 clf = DecisionTreeClassifier(criterion='gini',max_depth=15,min_samples_leaf=1,min_samples_split=3,splitter='best')
3 # Set a random state for reproducibility
4
5 # Train the classifier on the training data
6 clf.fit(x_train, y_train)
7
8 # Make predictions on the testing data
9 y_pred1 = clf.predict(x_test)
10
```

```
1 accuracy = accuracy_score(y_test, y_pred)
2
3
4 print("Accuracy:", accuracy)
5
```

Accuracy: 0.9688

By using Decision Tree model we got accuracy of 96.8%

✓ Using CNN

A CNN (Convolutional Neural Network) is kind of like that detective for images. It analyzes pictures by looking for tiny details, like edges, shapes, and colors.

```

1 #Loading the Data
2 img_rows, img_cols = 28, 28
3
4 x_train = x_train.reshape(60000, 28, 28, 1)
5 x_test = x_test.reshape(10000, 28, 28, 1)
6 # input_shape = (img_rows, img_cols, 1)
7
8 # print('input_shape: ', input_shape)
9 print('x_train shape:', x_train.shape)
10 print('x_test shape:', x_test.shape)
11 y_train.shape

x_train shape: (60000, 28, 28, 1)
x_test shape: (10000, 28, 28, 1)
(60000,)

1
2 # rescale to have values within 0 - 1 range [0,255] --> [0,1]
3 x_train = x_train.astype('float32')/255
4 x_test = x_test.astype('float32')/255
5
6 print('x_train shape:', x_train.shape)
7 print(x_train.shape[0], 'train samples')
8 print(x_test.shape[0], 'test samples')

x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples

1 from tensorflow.keras.models import Sequential
2 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
3
4 # build the model object
5 model = Sequential()
6
7 # CONV_1: add CONV layer with RELU activation and depth = 32 kernels
8 model.add(Conv2D(32, kernel_size=(3, 3), padding='same', activation='relu', input_shape=(28,28,1)))# 'same' in the sense it'll give the shape same as input
9 # POOL_1: downsample the image to choose the best features
10 model.add(MaxPooling2D(pool_size=(2, 2)))
11
12 # CONV_2: here we increase the depth to 64
13 model.add(Conv2D(64, (3, 3), padding='same', activation='relu'))
14 # POOL_2: more downsampling
15 model.add(MaxPooling2D(pool_size=(2, 2)))
16
17 # flatten since too many dimensions, we only want a classification output
18 model.add(Flatten())
19
20 # FC_1: fully connected to get all relevant data
21 model.add(Dense(64, activation='relu'))
22
23 # FC_2: output a softmax to squash the matrix into output probabilities for the 10 classes
24 model.add(Dense(10, activation='softmax'))
25
26 model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 28, 28, 32)	320
max_pooling2d (MaxPooling2D)	(None, 14, 14, 32)	0
conv2d_1 (Conv2D)	(None, 14, 14, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 64)	200768
dense_1 (Dense)	(None, 10)	650
=====		
Total params: 220234 (860.29 KB)		
Trainable params: 220234 (860.29 KB)		
Non-trainable params: 0 (0.00 Byte)		

```

1 # compile the model
2 model.compile(loss='categorical_crossentropy', optimizer='rmsprop',
3               metrics=['accuracy'])

1 from tensorflow.keras.utils import to_categorical
2 num_classes = 10
3 # print first ten (integer-valued) training labels
4 print('Integer-valued labels:')
5 # print(y_train[:10])
6
7 # one-hot encode the labels
8 # convert class vectors to binary class matrices
9 y_train = to_categorical(y_train, num_classes)
10 y_test = to_categorical(y_test, num_classes)
11

Integer-valued labels:

1 from tensorflow.keras.callbacks import ModelCheckpoint
2
3 # train the model
4 checkpointer = ModelCheckpoint(filepath='model.weights.best.hdf5', verbose=1,
5                               save_best_only=True)
6 hist = model.fit(x_train, y_train, batch_size=64, epochs=1,
7                 validation_data=(x_test, y_test), callbacks=[checker], verbose=2, shuffle=True)
8

Epoch 1: val_loss improved from inf to 0.07368, saving model to model.weights.best.hdf5
938/938 - 67s - loss: 0.1611 - accuracy: 0.9511 - val_loss: 0.0737 - val_accuracy: 0.9764 - 67s/epoch - 71ms/step

1 score = model.evaluate(x_test, y_test, verbose=0)
2 score
3

[0.07367992401123047, 0.9764000177383423]

```

By using CNN we got accuracy of 97.64%

▼ Summary:

The Summary here we give is as follows;

- We trained the model using KNN keeping the neighbourhood points as 5. We got accuracy of 97%.
- We Trained the model using decission Tree keeping criterion as gini,max depth of 15,minimum sample leaves as 1,minimum samples as 3 and splitter as best. We got accuracy of 96.8%.
- We trained the model using SVM keeping kernel as rbf since we are using image arrays, it'll be working well and gamma as auto. We got accuracy of 97.64%. Handwritten digit recognition is a challenging yet pivotal task in the realm of machine learning and computer vision. Various algorithms, such as Support Vector Machines (SVM), Random Forest, Gradient Boosting, k-Nearest Neighbors (k-NN), and Convolutional Neural Networks (CNN), have been employed to tackle this problem, each with its unique strengths and characteristics.

SVM demonstrates its efficacy in separating classes by finding the optimal hyperplane, making it a robust choice for linearly separable datasets. This algorithm has showcased remarkable success, achieving an impressive 98% accuracy.

Random Forest model excel in handling high-dimensional data, capturing complex patterns, and reducing overfitting by leveraging ensemble methods. This algorithm has showcased remarkable success, achieving an impressive 97% accuracy.

k-NN, with its simplicity and intuitive approach, identifies patterns based on similarity measures, making it effective for smaller datasets. However, its performance might be impacted by the curse of dimensionality for larger feature spaces. This algorithm has showcased remarkable success, achieving an impressive 97% accuracy.

On the other hand, CNNs have gained immense popularity and achieved remarkable success in image recognition tasks, especially for handwritten digit recognition. Their ability to automatically learn hierarchical representations of features from raw pixel data through convolutional layers makes them highly suitable for capturing spatial dependencies and intricate patterns within images. This algorithm has showcased remarkable success, achieving an impressive 97% accuracy.

- The network begins with a sequence of two convolutional layers, followed by max pooling layers.
- The final layer has one entry for each object class in the dataset, and has a softmax activation function, so that it returns probabilities.

- The Conv2D depth increases from the input layer of 1 to 32 to 64.
- We also want to decrease the height and width - This is where maxpooling comes in. Notice that the image dimensions decrease from 28 to 14 after the pooling layer.
- You can see that every output shape has **None** in place of the batch-size. This is so as to facilitate changing of batch size at runtime.
- Finally, we add one or more fully connected layers to determine what object is contained in the image.

Conclusion

- We conclude that the best machine learning model for the given dataset is KNN and Decision Tree of accuracy 97% & 96.8%. Comparing